

## **PROJECT NAME:**

### Linux Character Device Driver for System Metrics

## **INTRODUCTION**

This project involves developing a Linux character device driver to retrieve and output system metrics such as CPU usage, memory usage, and disk I/O metrics. The device driver can be loaded and unloaded from the kernel and provides a character device interface for user interaction.

## **Source Code:**

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/fs.h>
#include <linux/uaccess.h>
#include <linux/slab.h>
#include <linux/proc_fs.h>
#include <linux/timer.h>
#include <linux/sched.h>
#include <linux/jiffies.h>

#define DEVICE_NAME "sys_metrics" // Device name
#define CLASS_NAME "sys_class" // Device class name

static int majorNumber; // Major number for the device
```

```

static struct class* sysMetricsClass = NULL; // Device class pointer
static struct device* sysMetricsDevice = NULL; // Device pointer

static struct timer_list metrics_timer; // Timer for periodic metrics collection
static unsigned long interval = 5 * HZ; // Interval for the timer (5 seconds)

static char *metrics_buffer; // Buffer to store the metrics
static int buffer_size = 0; // Size of the buffer

// Function prototypes for device operations
static int dev_open(struct inode *inode, struct file *file);
static int dev_release(struct inode *inode, struct file *file);
static ssize_t dev_read(struct file *file, char *buffer, size_t len, loff_t *offset);
static ssize_t dev_write(struct file *file, const char *buffer, size_t len, loff_t *offset);

// File operations structure
static struct file_operations fops = {
    .open = dev_open,
    .read = dev_read,
    .write = dev_write,
    .release = dev_release,
};

// Device open function
static int dev_open(struct inode *inode, struct file *file) {
    printk(KERN_INFO "sys_metrics: Device opened\n");
    return 0;
}

// Device release function

```

```

static int dev_release(struct inode *inodep, struct file *filep) {
    printk(KERN_INFO "sys_metrics: Device closed\n");
    return 0;
}

// Device read function
static ssize_t dev_read(struct file *filep, char *buffer, size_t len, loff_t *offset) {
    int error_count = 0;

    // Check if the buffer is large enough
    if (len < buffer_size) return -EFAULT;

    // Copy the metrics data to user space
    error_count = copy_to_user(buffer, metrics_buffer, buffer_size);

    if (error_count == 0) {
        printk(KERN_INFO "sys_metrics: Sent %d characters to the user\n", buffer_size);
        return buffer_size;
    } else {
        printk(KERN_INFO "sys_metrics: Failed to send %d characters to the user\n",
error_count);
        return -EFAULT;
    }
}

// Device write function (not supported)
static ssize_t dev_write(struct file *filep, const char *buffer, size_t len, loff_t *offset) {
    printk(KERN_ALERT "sys_metrics: Write operation not supported\n");
    return -EINVAL;
}

// Function to collect system metrics

```

```

void get_system_metrics(void) {
    struct sysinfo sys_info;

    int num_cpus = num_online_cpus(); // Number of CPU cores

    int i;

    // Get memory information
    si_meminfo(&sys_info);

    // Format the metrics data into the buffer
    buffer_size = sprintf(metrics_buffer,
        "CPU cores: %d\n"
        "Total RAM: %lu kB\n"
        "Free RAM: %lu kB\n"
        "Shared RAM: %lu kB\n"
        "Buffered RAM: %lu kB\n"
        "Total Swap: %lu kB\n"
        "Free Swap: %lu kB\n"
        "Uptime: %lu seconds\n",
        num_cpus,
        sys_info.totalram * 4,
        sys_info.freeram * 4,
        sys_info.sharedram * 4,
        sys_info.bufferram * 4,
        sys_info.totalswap * 4,
        sys_info.freeswap * 4,
        sys_info.uptime);

    // Append CPU idle time for each core (this part needs proper implementation)

```

```

for (i = 0; i < num_cpus; i++) {
    buffer_size += sprintf(metrics_buffer + buffer_size,
                           "CPU %d: %lu\n",
                           i,
                           get_cpu_idle_time_us(i));
}
}

// Timer callback function
static void metrics_timer_callback(struct timer_list *timer) {
    printk(KERN_INFO "sys_metrics: Timer callback executed\n");
    get_system_metrics();
    mod_timer(&metrics_timer, jiffies + interval); // Restart the timer
}

// Module initialization function
static int __init sys_metrics_init(void) {
    printk(KERN_INFO "sys_metrics: Initializing the sys_metrics LKM\n");

    // Register the character device driver
    majorNumber = register_chrdev(0, DEVICE_NAME, &fops);
    if (majorNumber < 0) {
        printk(KERN_ALERT "sys_metrics failed to register a major number\n");
        return majorNumber;
    }

    printk(KERN_INFO "sys_metrics: registered correctly with major number %d\n",
           majorNumber);

    // Register the device class
    sysMetricsClass = class_create(THIS_MODULE, CLASS_NAME);
    if (IS_ERR(sysMetricsClass)) {

```

```

    unregister_chrdev(majorNumber, DEVICE_NAME);

    printk(KERN_ALERT "Failed to register device class\n");

    return PTR_ERR(sysMetricsClass);
}

printk(KERN_INFO "sys_metrics: device class registered correctly\n");


// Create the device

sysMetricsDevice = device_create(sysMetricsClass, NULL, MKDEV(majorNumber, 0),
NULL, DEVICE_NAME);

if (IS_ERR(sysMetricsDevice)) {
    class_destroy(sysMetricsClass);

    unregister_chrdev(majorNumber, DEVICE_NAME);

    printk(KERN_ALERT "Failed to create the device\n");

    return PTR_ERR(sysMetricsDevice);
}

printk(KERN_INFO "sys_metrics: device created correctly\n");


// Allocate memory for the metrics buffer

metrics_buffer = kmalloc(1024, GFP_KERNEL);

if (!metrics_buffer) {
    device_destroy(sysMetricsClass, MKDEV(majorNumber, 0));

    class_destroy(sysMetricsClass);

    unregister_chrdev(majorNumber, DEVICE_NAME);

    printk(KERN_ALERT "sys_metrics: Failed to allocate memory\n");

    return -ENOMEM;
}


// Setup and start the timer

timer_setup(&metrics_timer, metrics_timer_callback, 0);

```

```

mod_timer(&metrics_timer, jiffies + interval);

printk(KERN_INFO "sys_metrics: Device driver loaded successfully\n");
return 0;
}

// Module exit function
static void __exit sys_metrics_exit(void) {
    del_timer(&metrics_timer); // Delete the timer
    kfree(metrics_buffer); // Free the buffer memory
    device_destroy(sysMetricsClass, MKDEV(majorNumber, 0)); // Remove the device
    class_unregister(sysMetricsClass); // Unregister the device class
    class_destroy(sysMetricsClass); // Destroy the device class
    unregister_chrdev(majorNumber, DEVICE_NAME); // Unregister the character device driver
    printk(KERN_INFO "sys_metrics: Device driver unloaded successfully\n");
}

// Specify the module initialization and exit functions
module_init(sys_metrics_init);
module_exit(sys_metrics_exit);

MODULE_LICENSE("GPL");
MODULE_AUTHOR("SOURAV");
MODULE_DESCRIPTION("A Linux character device driver to output system metrics.");
MODULE_VERSION("1.0");

```

## **Makefile:**

```

obj-m += sys_metrics.o

all:

```

```
make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
```

clean:

```
make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

## **Functional Requirements:**

### **Device Driver:**

- Load and unload from the kernel.
- Use kernel APIs to retrieve metrics.
- IOCTL for CPU core averages, memory, and disk usage.

### **System Metrics:**

- Retrieve CPU, memory, and disk I/O metrics.
- Accurate and efficient data collection.

## **System Metrics Details**

### **Data Sources:**

- /proc filesystem for metrics.
- Define data format and structure for system info.

### **Metrics:**

- CPU usage (by core).
- Memory usage.
- Disk I/O.

## **Character Device Interface:**

### **Implementation:**

- Readable character device for outputting metrics.



- User space interface for reading metrics data.

## **User Interface:**

### **Commands:**

- Load device driver.
- Unload device driver.
- Read metrics by reading from the character device.

## **DOCUMENTATION:**

### **Code Comments and Explanations:**

- The source code contains inline comments explaining the purpose of each function, variable, and key section of code.
- Key components include the `get_system_metrics` function for collecting system metrics, the `metrics_timer_callback` for handling periodic metric collection, and standard file operations (`dev_open`, `dev_read`, `dev_write`, `dev_release`).

## **User Manual:**

### **Compiling the Device Driver:**

1. Ensure you have the necessary kernel headers installed:

```
sudo apt-get install linux-headers-$(uname -r)
```

2. Use the provided makefile to compile the driver:

```
make
```

### **Loading the Device Driver:**

1. Load the module into the kernel:

```
sudo insmod sys_metrics.ko
```

2. Verify the device file has been created:

```
ls /dev/sys_metrics
```

### Reading System Metrics:

1. Use the cat command to read from the device file:

```
cat /dev/sys_metrics
```

### Unloading the Device Driver:

1. Remove the module from the kernel:

```
sudo rmmod sys_metrics
```

2. Verify the device file has been deleted:

```
ls /dev/sys_metrics
```

## **Description of System Metrics Retrieval Methods:**

System metrics are retrieved using various kernel APIs:

- `si_meminfo(&sys_info)` is used to gather memory information.
- `get_cpu_idle_time_us(cpu_id)` (or similar function) is used to gather CPU usage information.

The `metrics_timer_callback` function periodically collects and updates these metrics.